

Introduction to Deep Learning

(Ensuring Full Literacy Training Workshop)

Muhammad Abdul-Mageed & Chiyu Zhang
muhammad.mageed@ubc.ca & chiyuzh@mail.ubc.ca
Deep Learning & NLP Lab

The University of British Columbia

February 17, 2021

Instructors . . .



Chiyu



Muhammad

Agenda

1 Overview of AI

2 Deep Learning Inspiration & Architectures

- Biological Inspiration
- Recurrent Neural Networks
- Convolutional Neural Networks

What is Artificial Intelligence?

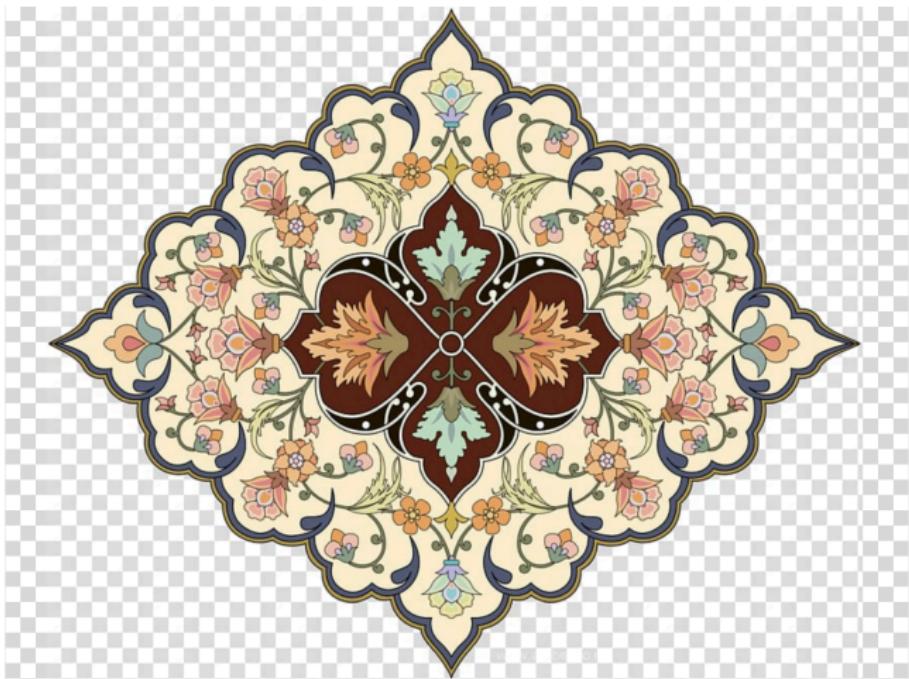


Art ...



Figure: Legendary inventors in Greek mythology. Pygmalion and the Statue [Wikimedia commons].

Art Contd.



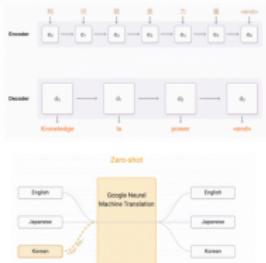
Can programmable machines become *intelligent*?



Figure: Ada Lovelace. 1842 [Wikipedia]

Deep Learning & Its Impact

Machine Translation

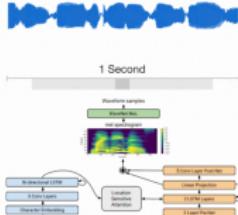
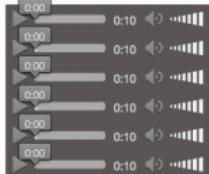


<https://research.googleblog.com/>



<https://translator.microsoft.com/help/articles/neural/>

Speech



<https://research.googleblog.com/>

Image (Captioning)



[Microsoft Research: He et al., 2015]



<https://research.googleblog.com/>

Overview

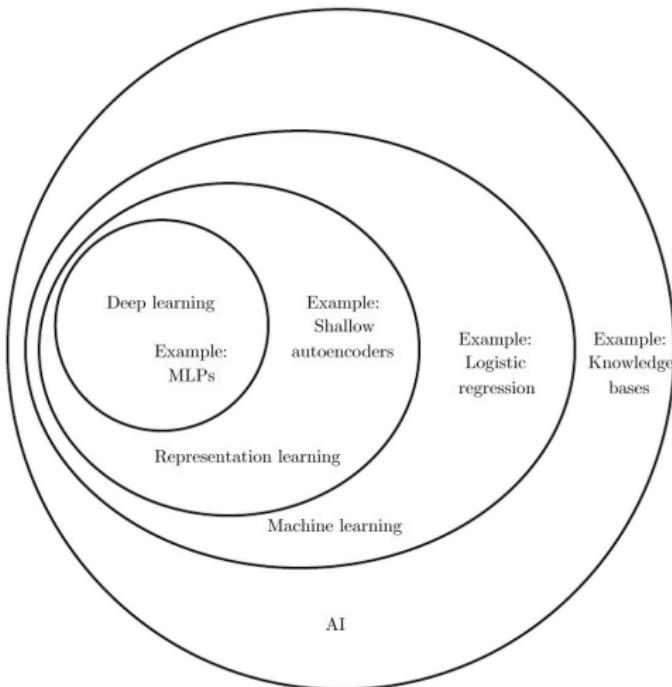


Figure: A Venn diagram situating DL within AI [Goodfellow et al., 2016].

Biological Inspiration

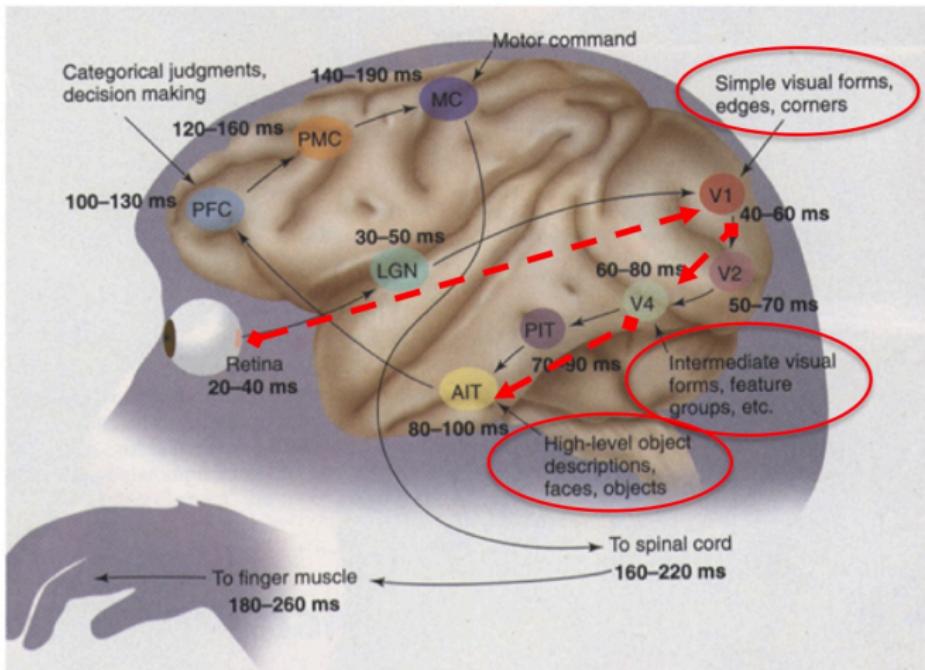


Figure: Visual cortex. [Modified from Simon Thorpe].

Networks with Multiple Layers

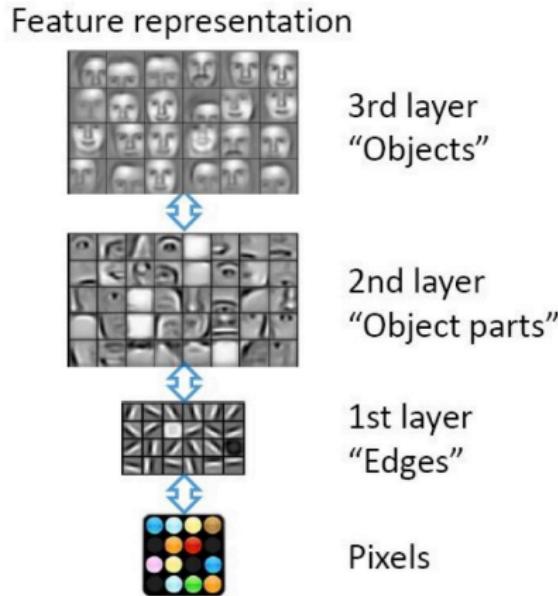


Figure: Learning in stages

Information processing in a neuron

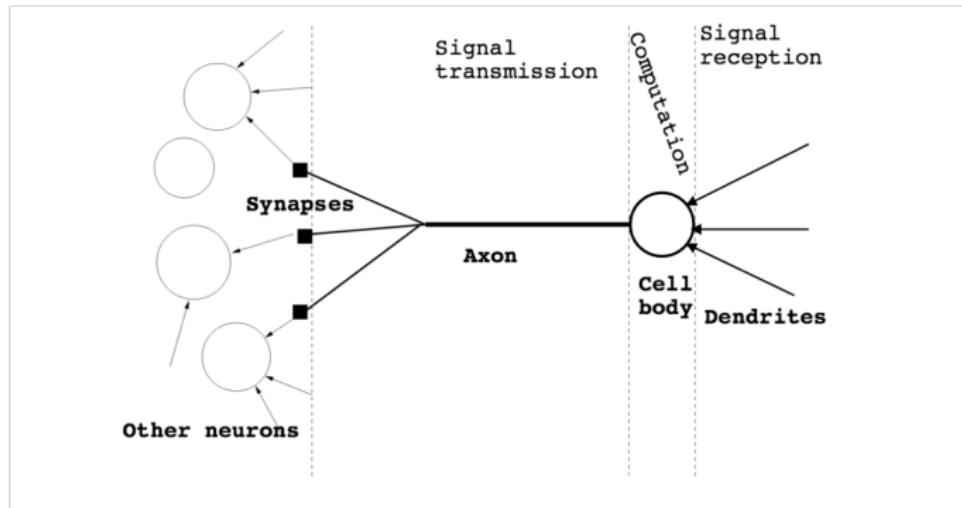


Figure: Information processing in a neuron. [Hyvarinen, Hurri, & Hoyer, 2009]

Neuron With Input

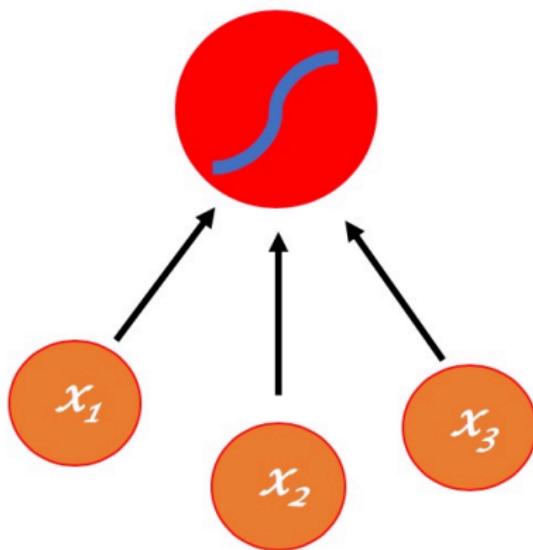


Figure: Input is a vector of \mathbf{x} with three units, each taking an index j

A Vector For Connection Weights

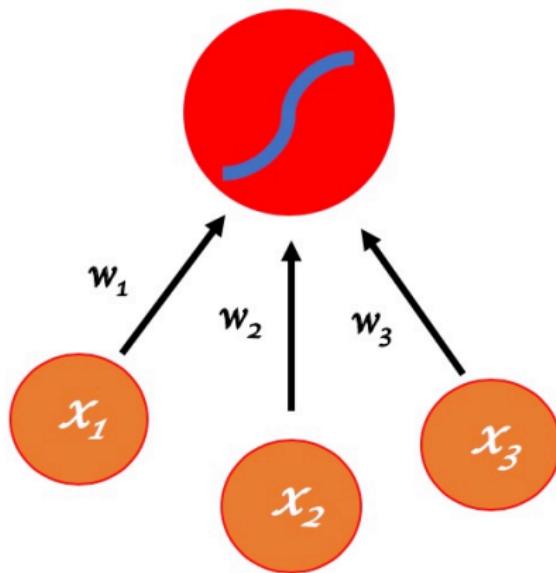


Figure: A vector of free parameters w with several items, each taking an index i

A Bias Unit

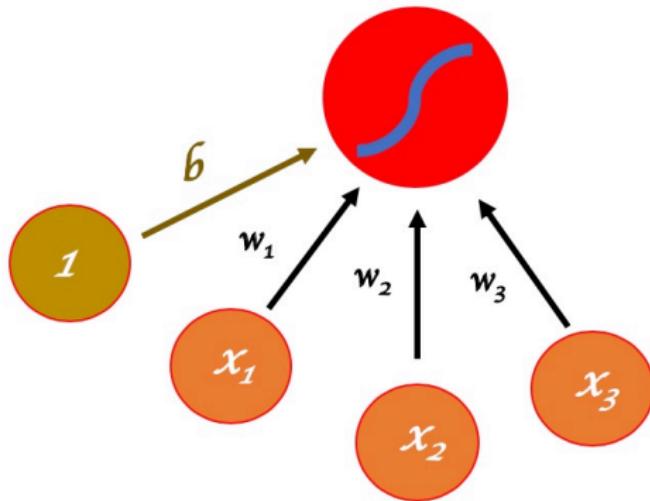


Figure: A bias unit b , equal to 1

A Deep Neural Network

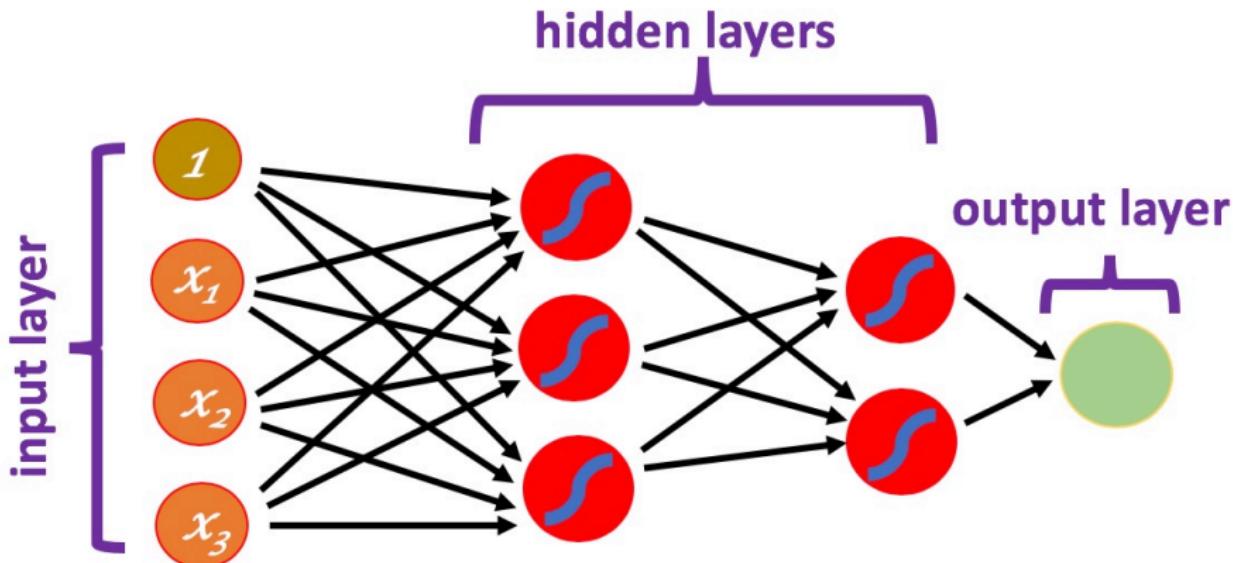


Figure: A **deep neural network**. The network has **2 hidden layers**.

RNN With Output

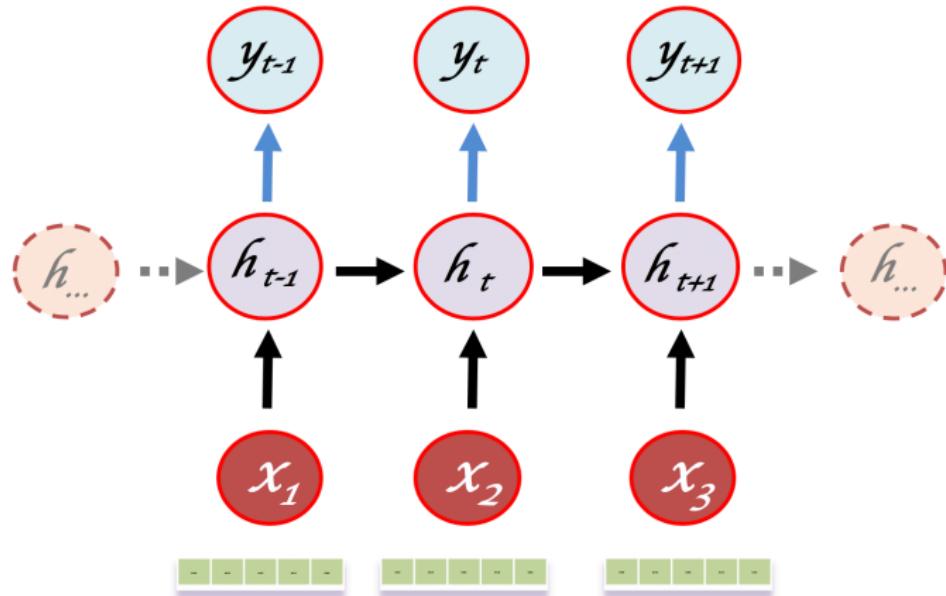


Figure: RNN is a generative model (can output a y at each time step).

Bidirectional RNNs

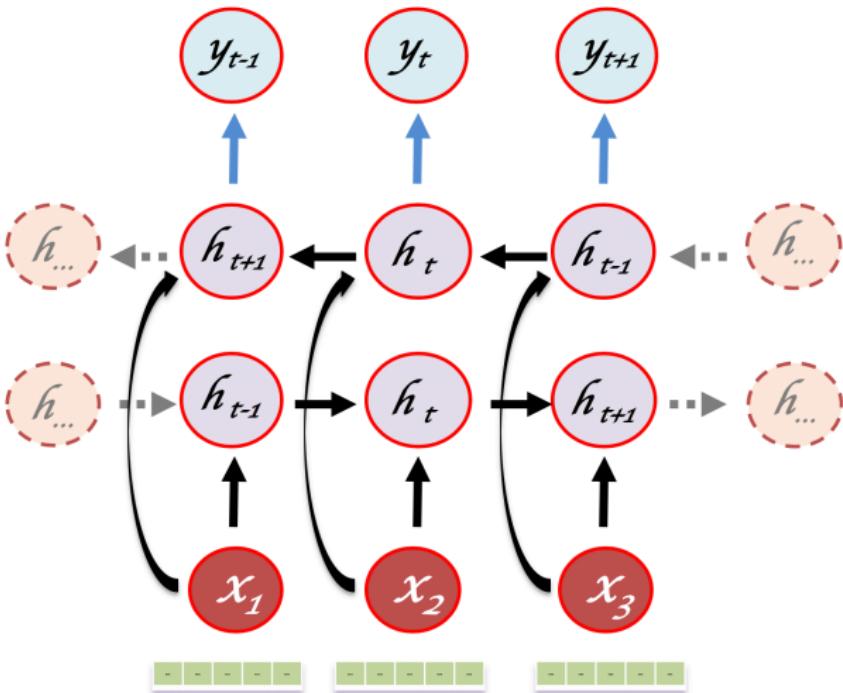


Figure: A bidirectional RNN (BiRNN).

Seq2Seq Models

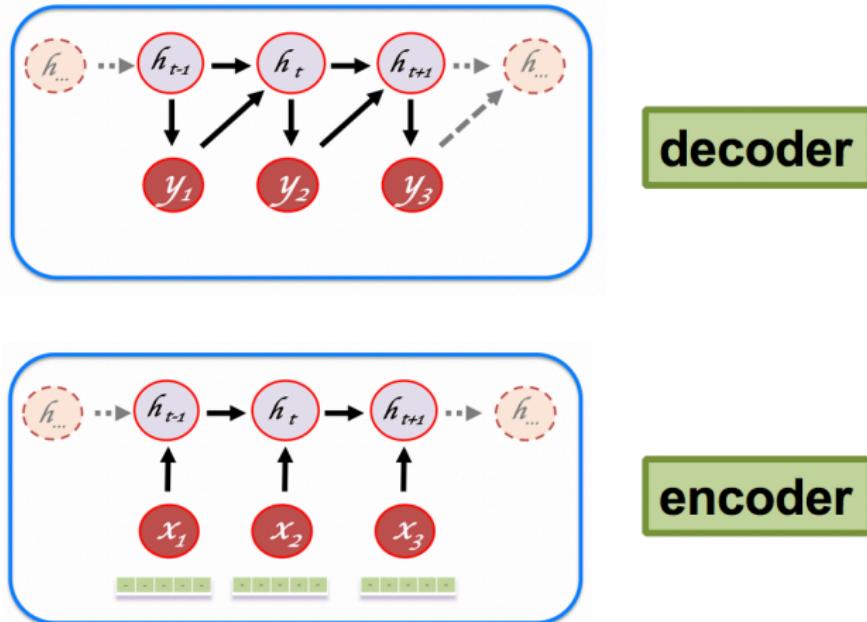


Figure: An Encoder and decoder.

Seq2Seq Models has Encoder Context

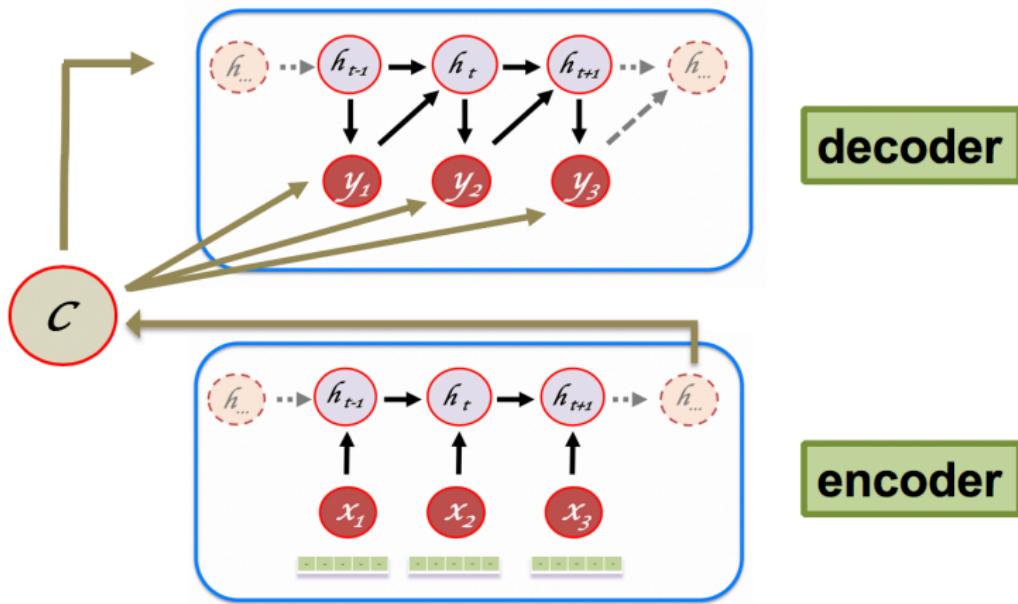
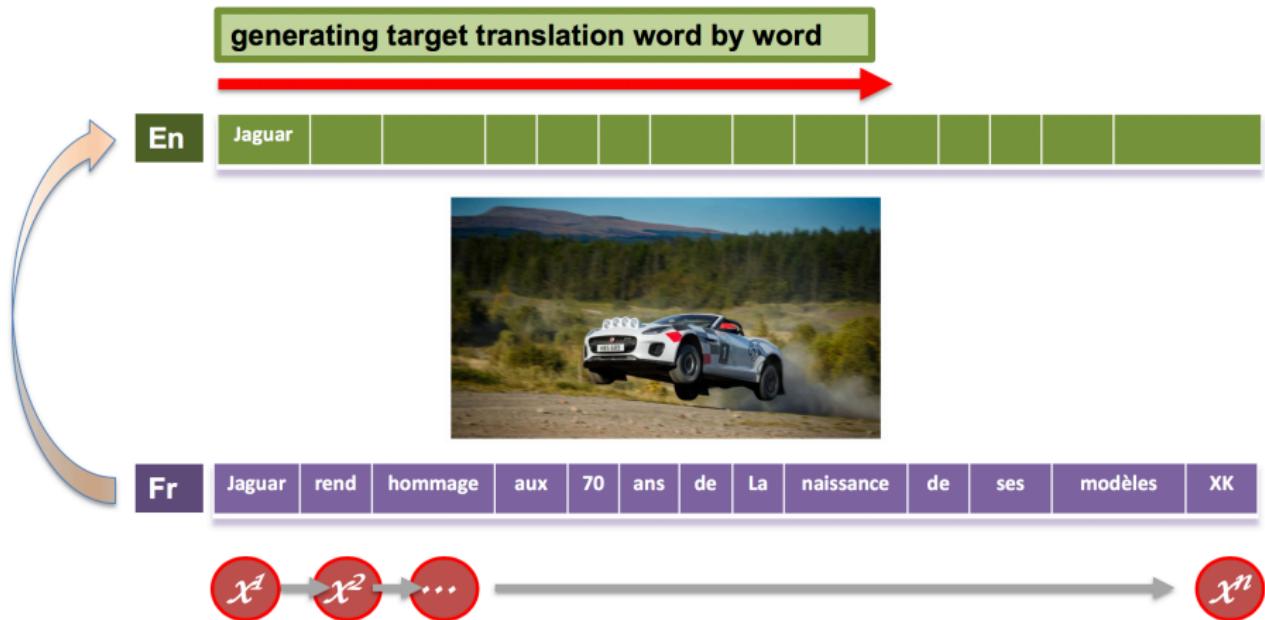
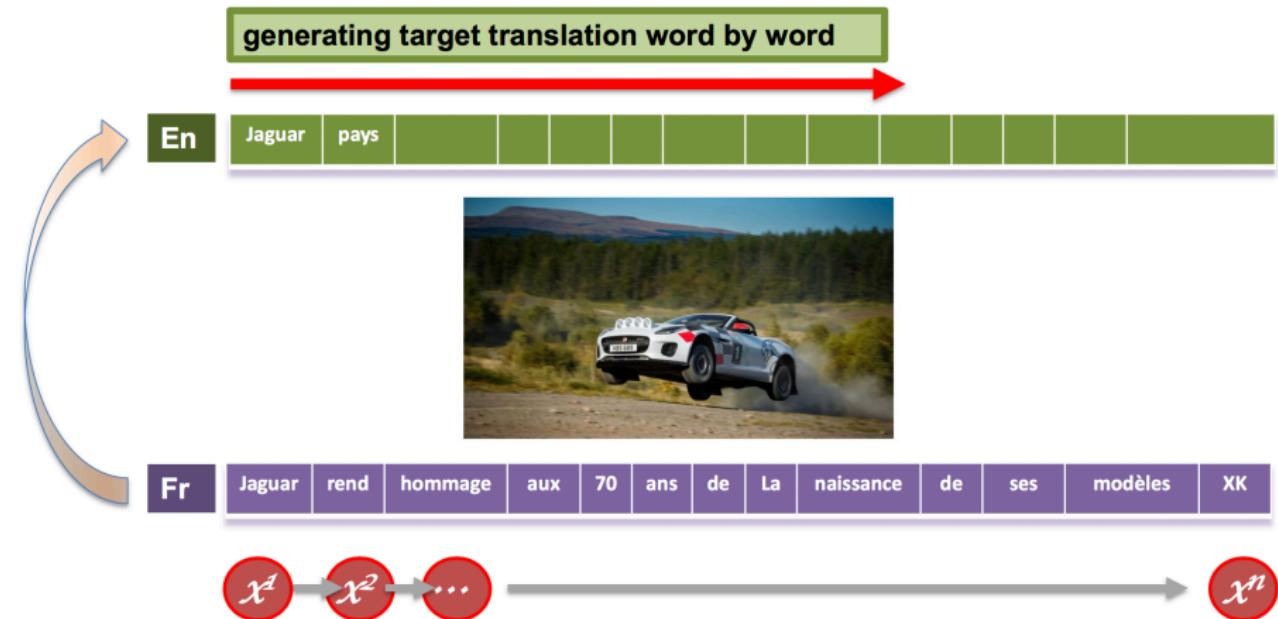


Figure: An Encoder and decoder connected via a **context vector C** .

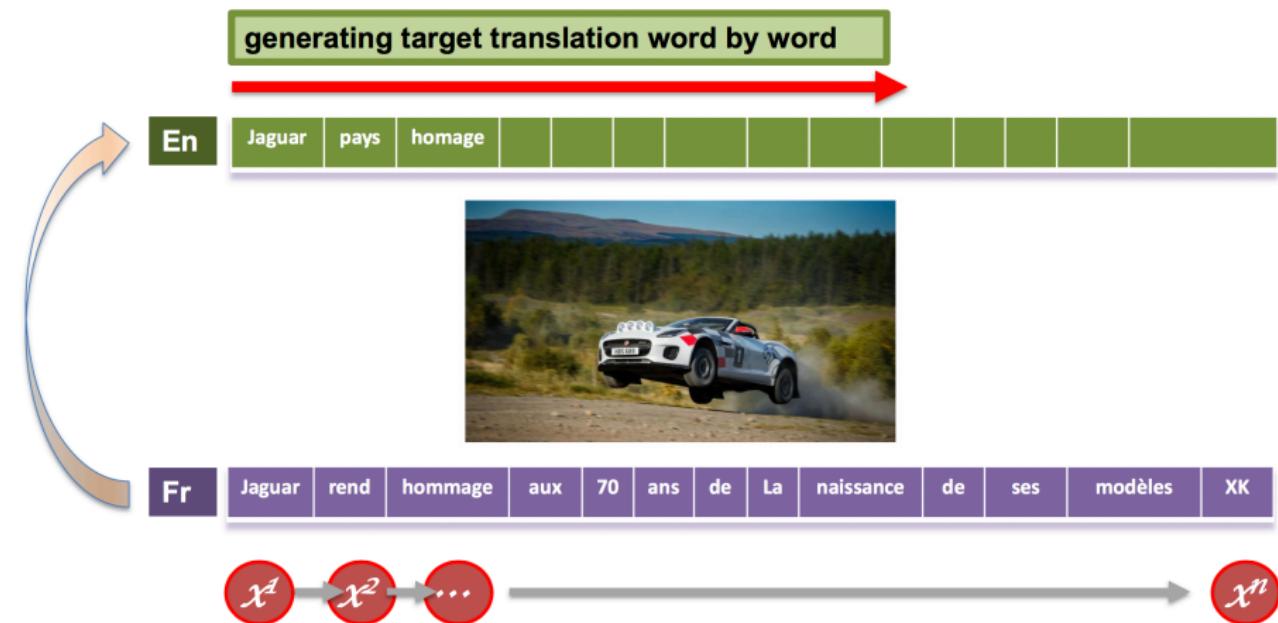
Generating Target (En) One Word at a Time



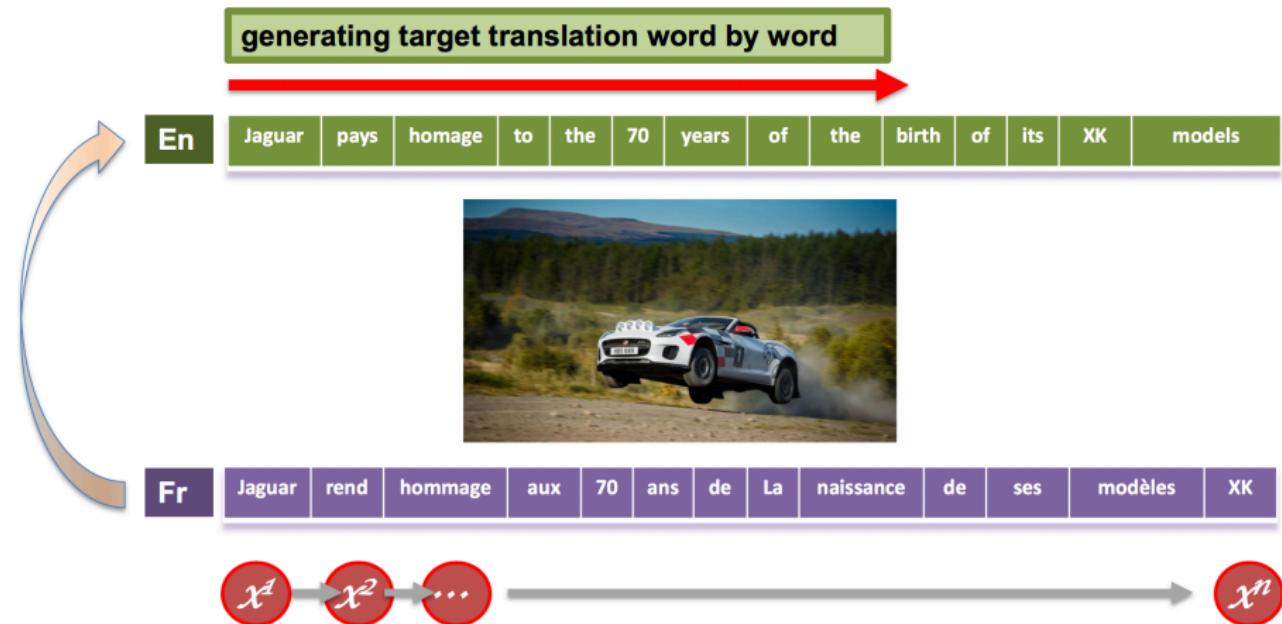
Generating Target (En) One Word at a Time *Contd. I*



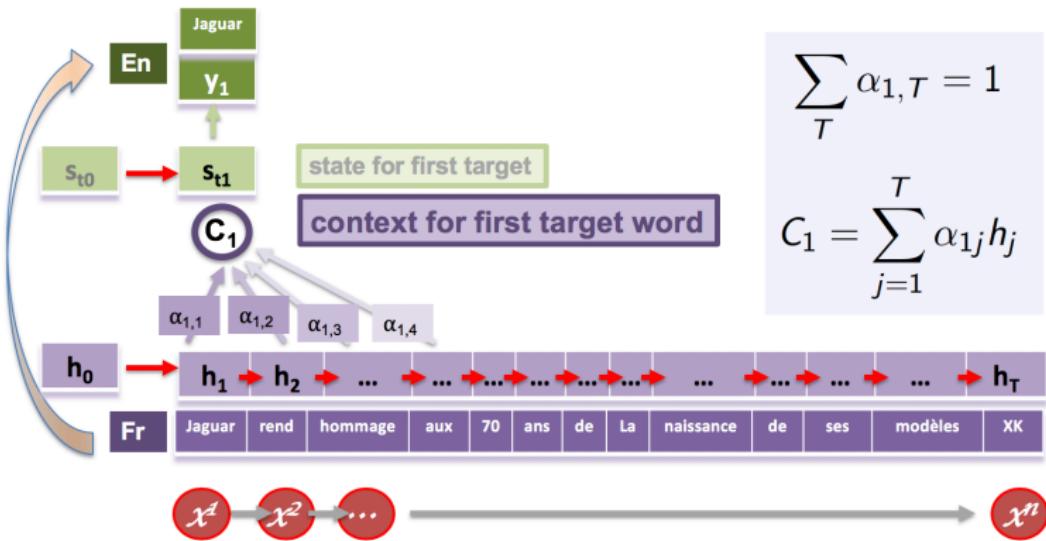
Generating Target (En) One Word at a Time *Contd. II*



Generating Target (En) One Word at a Time *Contd. III*



Context For State One



Context Ingredients

- α : attention weights. (We will learn these weights)
- h : activation from source

Convolutional Neural Networks

5 x 5 matrix (e.g., input image)

1	1	0	0	0
0	0	1	1	1
0	0	0	0	0
1	1	0	0	1
1	1	0	1	1

3 x 3 kernel (aka filter)

1	0	0
0	1	1
1	0	0

Desired feature map

Figure: **Input:** 5 x 5 matrix/2D Tensor. We will convolve with a 3 x 3 **kernel**, with a **stride** of 1. Our goal is to acquire an **feature map** (filling in the 3 x 3 matrix on right-hand side).

Convolving 1

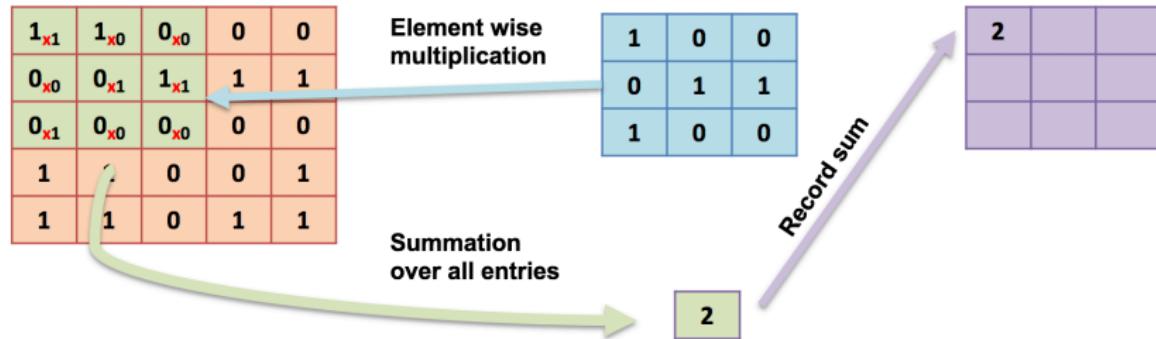


Figure: Filling in the first cell.

Convolving 2

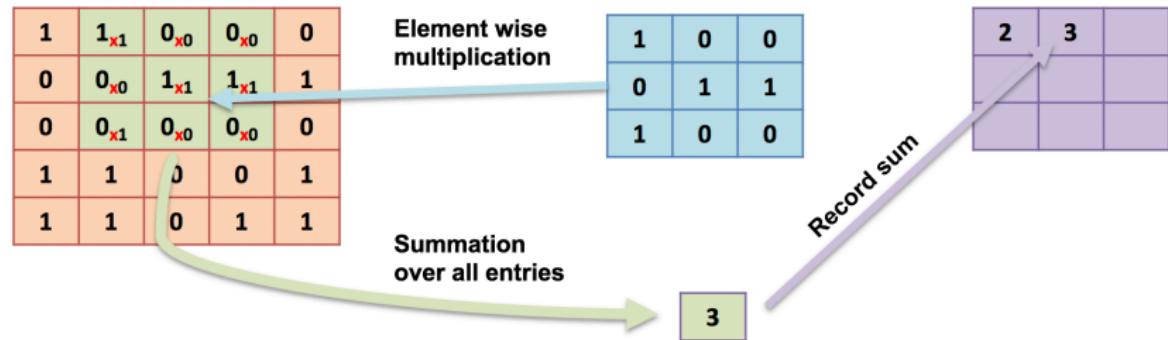


Figure: Filling in the second cell.

Convolving 3

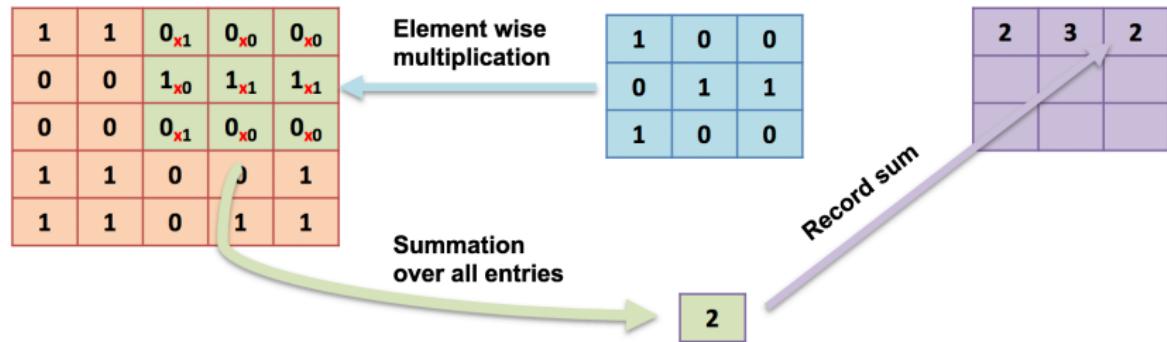


Figure: Filling in the third cell.

Convolving 4

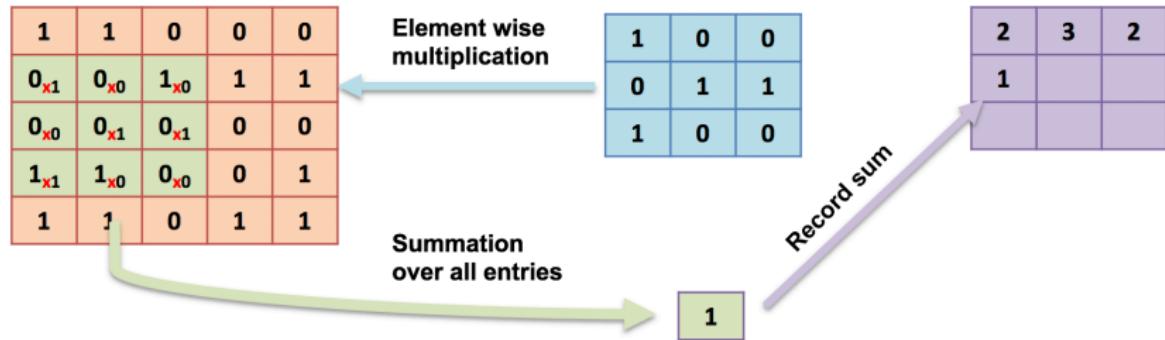


Figure: Filling in the fourth cell.

Convolving 5-9

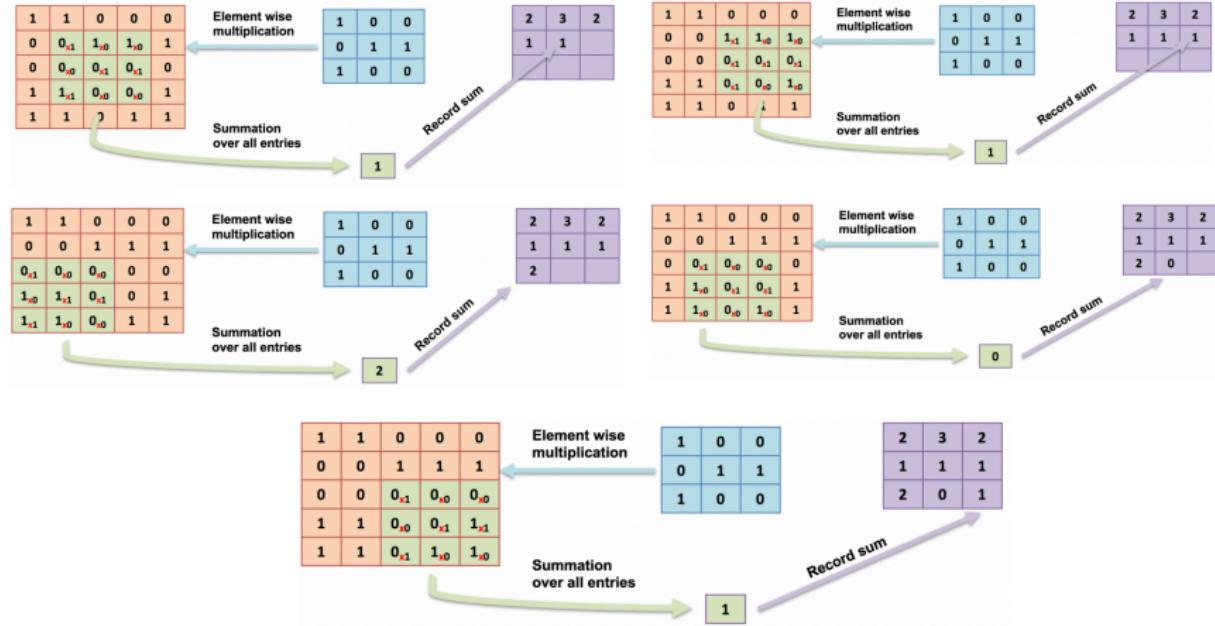


Figure: Filling in the fifth-to-ninth cells.

Convolving With 10 Kernels

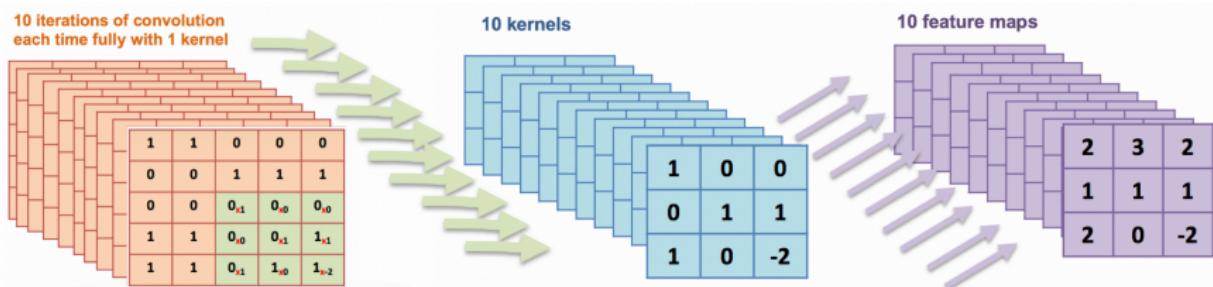


Figure: Convolving with 10 kernels, each with different values, we acquire 10 different feature maps.

Non-Linearity

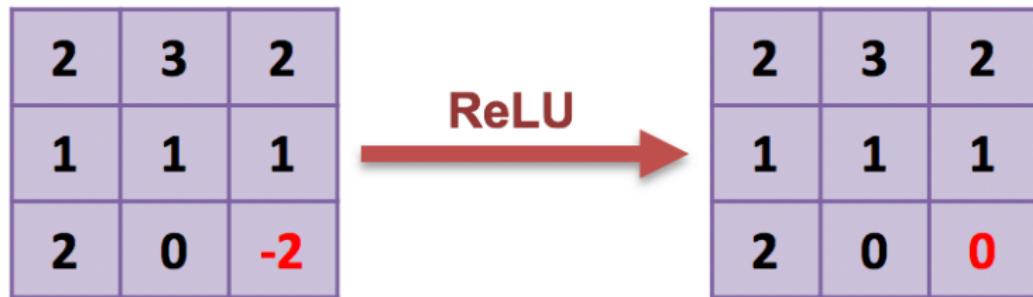


Figure: We actually apply a non-linear function like ReLU on the output of convolution operations such that we end up with meaningful feature maps.

Convolution With Non-Linearity

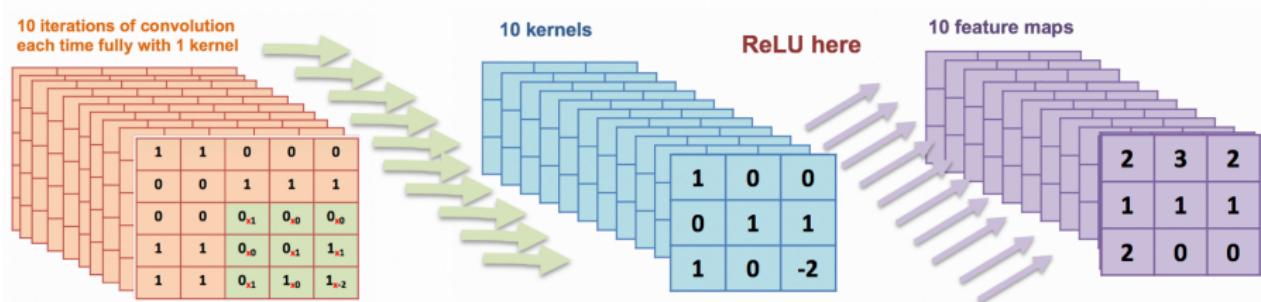


Figure: We actually apply a non-linear function like ReLU on the output of convolution operations such that we end up with meaningful feature maps.

Thank You!

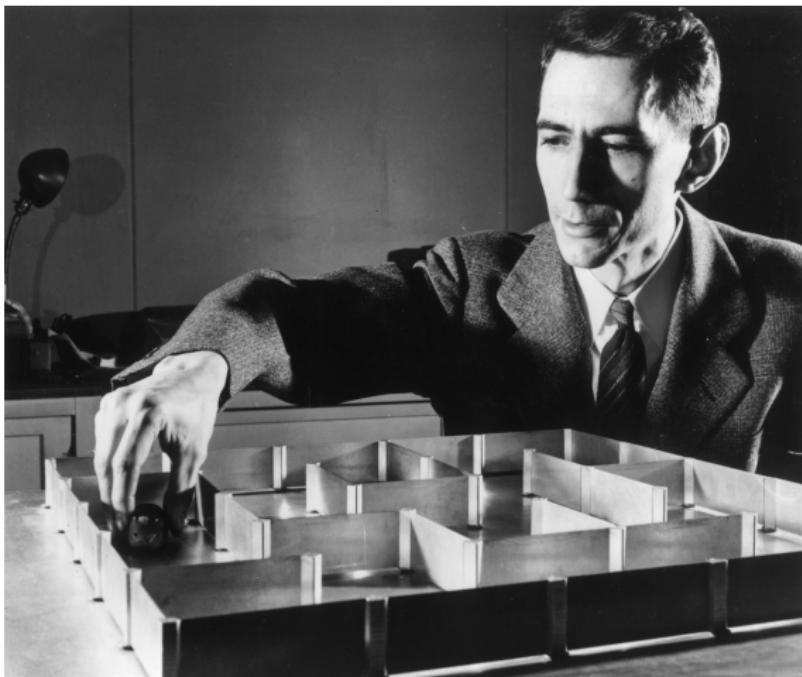


Figure: Claude Shannon. [From Time]